

Student ID: 700743212

Name: Gundlapally sreenidhi

NN & DL - Assignment 3

Video Link:

https://drive.google.com/file/d/1n_nYTv_sZ_YpEEXujkyAo9BNwy1hAtGv/view?usp=s_haring

GitHub Link: <https://github.com/sreenidhi-ux/assignment3>

Question 1:

Importing set of libraries, training and testing the data

+ Code + Text

```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.constraints import maxnorm
from keras.optimizers import SGD
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.utils import np_utils
```

```
[ ] np.random.seed(7)
```

```
[ ] (X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 2s 0us/step

```
[ ] y_train = np_utils.to_categorical(y_train)
    y_test = np_utils.to_categorical(y_test)
    num_classes = y_test.shape[1]
```

```
[ ] model = Sequential()
    model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
    model.add(Dropout(0.2))
    model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
    model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
    model.add(Flatten())
    model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
    model.add(Dropout(0.5))
    model.add(Dense(num_classes, activation='softmax'))
```

```
[ ] sgd = SGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
    model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
    print(model.summary())
```

Model: "sequential"

```
sgd = SGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130
Total params: 4,210,090		
Trainable params: 4,210,090		
Non-trainable params: 0		
None		

Fitting the data and evaluating the scores for accuracy

```
[ ] epochs = 5
    batch_size = 32
    model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
```

```
Epoch 1/5
1563/1563 [=====] - 302s 192ms/step - loss: 1.6841 - accuracy: 0.3892 - val_loss: 1.3400 - val_accuracy: 0.5142
Epoch 2/5
1563/1563 [=====] - 316s 202ms/step - loss: 1.3216 - accuracy: 0.5272 - val_loss: 1.1922 - val_accuracy: 0.5684
Epoch 3/5
1563/1563 [=====] - 316s 202ms/step - loss: 1.1505 - accuracy: 0.5912 - val_loss: 1.1238 - val_accuracy: 0.6010
Epoch 4/5
1563/1563 [=====] - 310s 199ms/step - loss: 1.0075 - accuracy: 0.6454 - val_loss: 1.0171 - val_accuracy: 0.6423
Epoch 5/5
1563/1563 [=====] - 296s 189ms/step - loss: 0.8995 - accuracy: 0.6832 - val_loss: 0.9656 - val_accuracy: 0.6657
<keras.callbacks.History at 0x7f0b836071c0>
```

```
[ ] scores = model.evaluate(X_test, y_test, verbose=0)
    print("Accuracy: %.2f%%" % (scores[1]*100))
```

Accuracy: 66.57%

Accuracy we got here is: **66.57%**

Finding the performance change in the below code

```
[1] import numpy as np
    from keras.datasets import cifar10
    from keras.models import Sequential
    from keras.layers import Dense, Dropout, Flatten
    from keras.layers.convolutional import Conv2D, MaxPooling2D
    from keras.constraints import maxnorm
    from keras.utils import np_utils
    from keras.optimizers import SGD

    # Fixing the random seed for reproducibility
    np.random.seed(7)

    # Loading the data
    (X_train, y_train), (X_test, y_test) = cifar10.load_data()

    # Normalizing inputs from 0-255 to 0.0-1.0
    X_train = X_train.astype('float32') / 255.0
    X_test = X_test.astype('float32') / 255.0

    # One hot encode outputs
    y_train = np_utils.to_categorical(y_train)
    y_test = np_utils.to_categorical(y_test)
    num_classes = y_test.shape[1]
```

Creating

And compiling the model

```
# Creating the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

Fitting and evaluating the model.

```

▶ # Compiling model
epochs = 5
learn_rate = 0.01
decay_rate = learn_rate / epochs
sgd = SGD(lr=learn_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fitting the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluating the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

```

Output:

▶

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dropout_3 (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1024)	2098176
dropout_4 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

=====
 Total params: 2,915,114
 Trainable params: 2,915,114

```

Total params: 2,915,114
Trainable params: 2,915,114
Non-trainable params: 0

/usr/local/lib/python3.10/dist-packages/keras/optimizers/legacy/gradient_descent.py:114: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
None
Epoch 1/5
1563/1563 [=====] - 598s 381ms/step - loss: 1.9256 - accuracy: 0.2889 - val_loss: 1.6807 - val_accuracy: 0.3877
Epoch 2/5
1563/1563 [=====] - 598s 382ms/step - loss: 1.5471 - accuracy: 0.4336 - val_loss: 1.4688 - val_accuracy: 0.4670
Epoch 3/5
1563/1563 [=====] - 594s 380ms/step - loss: 1.4079 - accuracy: 0.4881 - val_loss: 1.3898 - val_accuracy: 0.4968
Epoch 4/5
1563/1563 [=====] - 584s 374ms/step - loss: 1.3209 - accuracy: 0.5222 - val_loss: 1.2574 - val_accuracy: 0.5492
Epoch 5/5
1563/1563 [=====] - 575s 368ms/step - loss: 1.2645 - accuracy: 0.5439 - val_loss: 1.2194 - val_accuracy: 0.5616
Accuracy: 56.16%

```

Accuracy we got here is: **56.16%**.

The performance of the model is likely to improve with the addition of more layers and higher number of feature maps, but it will also increase the complexity and the training time of the model. The new model architecture provided in the instruction includes several new layers and higher number of feature maps, which may improve the accuracy of the model.

Question 2:

Predicting the first 4 images of the test data

```

[13] # Predicting the first 4 images of the test data
predns = model.predict(X_test[:4])
# Converting the predictions to class labels
predicted_lbls = np.argmax(predns, axis=1)
# Converting the actual labels to class labels
actual_lbls = np.argmax(y_test[:4], axis=1)

# Printing the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_lbls)
print("Actual labels:  ", actual_lbls)

```

```

1/1 [=====] - 0s 113ms/step
Predicted labels: [3 8 8 8]
Actual labels:   [3 8 8 0]

```

Question 3:

Importing the libraries and plottinh the training and validation accuracy

```
import matplotlib.pyplot as plt

# Plotting the training and validation loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

# Plotting the training and validation accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='lower right')
plt.show()
```

Output:

